Sana Nasir

Generative AI Engineer | AI Evaluation Specialist | Agentic AI & LLM Systems | RAG Specialist

0345-2262992 | sadsunsuf@gmail.com | Karachi, Pakistan
LinkedIn: linkedin.com/in/sana-nasir-521937316 | Portfolio: sadsunsuf-lgtm.github.io/my-portfolio
GitHub: github.com/sadsunsuf-lgtm

PROFESSIONAL SUMMARY

Production-focused Generative AI Engineer and AI Evaluation Specialist with six months of hands-on experience designing, validating, and quality-controlling LLM evaluation tasks across scientific and technical subjects. Skilled at building unambiguous, difficult problems with gold-standard step-by-step solutions, identifying reasoning errors and logical inconsistencies in model outputs, and passing multi-stage QC and reviewer feedback cycles. Also experienced in building LLM systems, multi-agent pipelines, and RAG platforms, with 11 production-grade AI projects delivered at DevelopersHub Corporation (top-performance award). Proficient in LangGraph, LangChain, FastAPI, Django REST Framework, Next.js, and Docker. Open to remote AI evaluation roles and full-time AI/ML roles globally.

TECHNICAL SKILLS

AI Evaluation & Quality: LLM Evaluation, Benchmark & Problem Design, Gold-Standard Solution Writing, Reasoning-Error Analysis, Model Output Review, Well-Posedness & Correctness Auditing, Test-Case Design, QC & Reviewer Feedback Cycles, Technical Content Review, Annotation; Subject Areas: Physics, Chemistry, Materials Science, Computer Science, Data Science, Biomedical Sciences

GenAI & LLM: LangGraph (Multi-Agent), LangChain, RAG Systems, Prompt Engineering, Agentic Web-Intelligence, Chatbot Development, Conversational AI, Multi-turn Chat Systems, LLaMA 3.3, Gemini 1.5/2.5 Pro, Mistral, Groq LPU

Frameworks: FastAPI, Django 4.2 (SaaS Architecture), Django REST Framework, Next.js 16 (App Router), TypeScript, Streamlit, Gradio

Integrations: WhatsApp API, Discord Webhook, Pollinations AI, Web Speech API, Leaflet.js, REST API Design, Streamlit Chat UI

ML / DL: TensorFlow, Keras, PyTorch, CNN (Image Classification), RNN/LSTM, ANN, GANs, PCA, Data Augmentation, BERT/DistilBERT fine-tuning (HuggingFace), Transfer Learning, Neural Networks, Scikit-Learn Pipelines, Cross-validation, GridSearchCV, Model Evaluation (F1, ROC-AUC, Confusion Matrix), Feature Engineering, Data Cleaning & Preprocessing

Data Science: Statistical Analysis, Probability, EDA (Exploratory Data Analysis), Data Visualization, Hypothesis Testing, Correlation Analysis, Pandas, NumPy, Matplotlib, Seaborn

Vector & Data: FAISS (Vector DB), PostgreSQL, Redis, DuckDB, spaCy (Entity Extraction), scikit-learn, Pydantic v2, Tenacity, DuckDuckGo API

NLP: HuggingFace Transformers, Tokenization, Text Embeddings, Sentiment Analysis, Text Classification, Zero-Shot & Few-Shot Prompting, Word2Vec, NER

CV & OCR: OpenCV (Advanced Preprocessing), Tesseract OCR, Document Intelligence for auditing & resume extraction

DevOps & Cloud: Docker, Git, Vercel, Google Cloud Run, Streamlit Cloud, HuggingFace Spaces, Google Colab, Kaggle, Jupyter, CI/CD, Secrets Management

Engineering: Robust API Integration, GDPR/CCPA Compliance, Error Handling, JWT Auth, SMTP Automation, Microservices Architecture

Languages: Python (OOP, Scripting), TypeScript, JavaScript, SQL

EXPERIENCE

AI Evaluation Specialist (Contract, Remote) May 2026 – Oct 2026

- Designed challenging multi-step scientific and technical coding problems as structured Jupyter notebooks (subproblems plus a main task), built to expose reasoning failures in state-of-the-art LLMs.
- Wrote precise gold-standard solutions and validated each task for scientific correctness and well-posedness (a single unambiguous answer, a complete specification) before submission.
- Took tasks through multi-stage QC and reviewer feedback, debugging and redesigning them to meet model-difficulty thresholds, test-coverage, convergence, and API-consistency requirements.
- Evaluated model-generated responses and documented reasoning errors, logical inconsistencies, numerical issues, and technical limitations.
- Worked across multiple subject areas, including physics, chemistry, and materials science, adapting problem design to each domain.
- Contributed to agent-evaluation and task-hardening work on a simulated fintech platform using Docker, GCP, Cloud Run, and DuckDB.

AI & Machine Learning Intern — DevelopersHub Corporation Apr 2026 – May 2026 · Karachi

- Delivered 11 production-grade AI projects across 2 phases; earned top-performance award.
- **Phase 1 — ML & Data Science:** EDA, Linear/Logistic Regression, K-Means, Random Forest, hyperparameter tuning on real-world datasets.

- **Phase 2 — Agentic GenAI:** Agentic PDF RAG Chatbot (LangChain + FAISS + Groq/Docker); BERT fine-tuning for news classification (HuggingFace Spaces); Zero-Shot vs. Few-Shot prompt engineering; Multimodal Price Predictor; Customer Churn ML pipeline.

KEY PROJECTS

AI RecruitPro — Talent Management Platform | Django · React · spaCy · Gemini · Docker
- Enterprise-grade SaaS with spaCy resume parsing, sentiment-aware video analysis, automated bias detection, and predictive candidate ranking.

TalentGuard — Resume Screener | Python · Streamlit · Gemini 1.5 Flash · Tenacity
- Batch-processes PDFs against JDs; scores/ranks candidates; AI Executive Memo generator; exponential-backoff retry logic.

Multi-Agent Research System | FastAPI · LangChain · PostgreSQL · Docker · Redis
- 5-agent pipeline (Research, Analysis, Validation, Writer, Viz); 80% research-time reduction; JWT auth; cloud-ready.

AI Health Assistant Chatbot | Next.js 16 · TypeScript · Groq · LLaMA 3.3 · TailwindCSS
- High-speed conversational AI chatbot for health education; Groq LPU inference for sub-second responses; safety guardrails via prompt engineering; CI/CD deployed.

SEO_FORGE Pro (SaaS) | Gemini 2.5 · LLaMA 3.3 · Groq · LangChain
- Dual-inference 'Battle Mode' content engine; 98% workflow reduction; 500+ daily users at 99.8% uptime.

CIRO — Crisis Intelligence & Response Orchestrator | Google Antigravity · Gemini 1.5 Flash · Leaflet.js · WhatsApp API · Web Speech API
- Google Hackathon — Pakistan's first bilingual (Urdu/English) AI emergency system; 15 Antigravity agents fuse social media, weather & traffic to detect urban crises in under 2 minutes across 7 cities.

Bolta Bazaar AI | Google Cloud Run · Gemini · Whisper · Docker
- Voice-first inventory management system for local shopkeepers using Urdu audio-to-intent pipelines; fully containerized and deployed live on Google Cloud Run.

EDUCATION

SMIT (Saylani Mass IT Training) · AI & Data Science May 2025 – Dec 2025
University of Karachi · M.Sc. Physiology & Pathophysiology 2004 – 2007

CERTIFICATIONS

- Introduction to Modern AI — Anthropic (Badge & Certificate)
- Python Essentials 1 & 2 — Cisco NetAcad (Badge & Certificate)
- Introduction to Data Science — Cisco NetAcad (Badge & Certificate) · Issued Oct 2025
- AI Seekho 2026 — Google Antigravity National Hackathon — Certificate of Participation (Google for Developers, Tech Nation Pakistan, Telenor 4G)